


Agent Workforce Governance

PAST LEADERBOARD (PASSIVE)				
AGENT_ID	TASK_COUNT	LATENCY_MS	SUCCESS_RATE	STATUS
AGENT_001	1200	320	73%	PASSIVE
AGENT_002	1200	320	73%	PASSIVE
AGENT_003	1200	320	73%	PASSIVE
AGENT_004	1200	320	73%	PASSIVE
...	...	...	...	...


Fleets managed blind.



Fleets managed like a team.

Moving from passive LLM evaluation to active, evidence-based workforce governance.

The blind fleet leaves four critical operational questions unanswered




Existing tools hand you an evaluation score and stop. A number on a dashboard does not manage a workforce.

Current tools solve fractions of the problem in isolated silos

Capability	Eval Tools (LangSmith)	Routers (Portkey)	AI Governance (ServiceNow)	Agent Workforce Platform
Measures Quality?	✓	✗	✗	✓
Routes on measured competence?	✗	✓ (Price only)	✗	
Earns Autonomy?	✗	✗	✗	
Provides Audit?	✗	✗	✓	


The capability gap is not missing features—it is the lack of a decision layer connecting measured competence to governed action.

Tying validated competence to governed action




We do not replace your agents. We manage them based on evidence.

A dual-stream architecture separating measurement from operations



The Measurement Engine operates a deterministic multiagent judge



The Panel:
5 parallel LLM judges.


The Aggregator:
Deterministic logic combines scores and enforces a strict safety veto.

The Output:
1-5 scoring, pass/fail verdict, and written rationale.

Validated against human gold labels before driving automation


98.7%

Pass/Fail Accuracy




0.97

Cohen's Kappa
(Near-perfect human agreement)




0.98

Spearman Correlation



Trust must be proven. The measurement engine operates with established human-level reliability, clearing all baseline validation targets against a 76-item human-labelled dataset.

Preserving perfect accuracy at a fraction of the cost



The Problem:

A 5-judge panel running on flagship models is financially unscalable.


The Cascade Solution:

Screen items via the cheaper model first. Escalate only when confidence is low.

The Result:

100% pass/fail agreement preserved at $\sim 1/8$ th the cost of the full panel (\$0.0010 vs \$0.0080 per item).

Quality scores hide operational sins



True governance requires combining the final quality score with the operational cost of achieving it. An agent might eventually produce a perfect answer, but burn through tokens and retries to get there.

Generating thirty days of operational history on demand

Synthetic Mode



- Calibrated runtime simulator
- Injects specific operational incidents (latency spikes, degrading agents)
- Spans a 30-day timeline offline


Live Mode



- Real LangGraph agents
- Executes authentic retrieve/verify/retry loops
- Emits genuine telemetry traces


Provides credible runtime dynamics, drift scenarios, and tier changes for immediate governance testing without the cost of running live API calls for a month.

The Decision Layer translates history into automated oversight



Trust is earned from evidence, not assumed by default

The Autonomy Staircase




Agent: trans_weak

Status: Demoted to Blocked

Rationale: Despite sporadic quality passes, its operational error rate exceeded 20%, triggering a hard governance gate.


Routing to the most efficient agent that is demonstrably good enough



Decision: The router dynamically selects `rag_cheap`. Historical data proves it handles this task family with equal quality but superior operational speed, bypassing the slower model.

Detecting silent regressions and enforcing an accountable trail

Performance Over Time



ALERT

Alert Box


Incident: sum_weak performance drop.
Alert: groundedness_drift of -0.23 points detected.

Audit Trail

Timestamp	Agent	Trigger	Action Logged
[Current Time]	sum_weak	Drift Alert	Autonomy Demotion Enforced

The alert, the resulting autonomy demotion, and the exact timestamp are permanently written to the compliance log. The definitive record of who decided what, when, and why.

From measured competence to governed action



Cost & Efficiency

The Cascade Judge cuts evaluation costs by 7/8ths, while smart routing prevents overspending on overpowered models.



Quality & Safety

Hard gating mechanisms and real-time operational drift alerts catch regressions before they impact the user.



Definitive Accountability

A single, unified audit trail proving why every automated decision was made, bridging the gap between engineering and compliance.